

Announcements

- Homework 2 should be released soon. It will be due on Feb. 28th at 5pm CST.
- Lab 1 reports are due at the beginning of Lab 2 this week.
- Lab 2 starts this week.

Overview

- Last lecture: basic concepts in periodic tasks.
 - Real-time periodic tasks and the problems of jitter and drifts.
- Today: Implementation of real time periodic task in POSIX RT.
 - Timer signals (software interrupt)
 - Signal handlers (Interrupt Service Routines – ISR)

Review: Potential Timing Problems

- E1: Process could be swapped out of memory (causing drift and jitter).
- Solution: Pin it down in memory! In LynxOS, processes with priority greater than 16 (default) will not be swapped out.

```
1. current_time = read_clock()           //E2 if preempted, drift
2. If (START_TIME - current_time < 10 msec) { //report too late and exit}
3. sleep(START_TIME - current_time)      //E3: if preempted, drift
loop
    4. current_time = read_clock()         // same problem as line 1
    5. wake_up_time = current_time + 20 msec //E4: drift if current time is delayed at line 4
    6. Read sensor data                   //E5: if preempted, input jitter
    7. //do work
    8. current_time = read_clock()         // same problem as line 1
    9. //send control data to the device   //E6: output jitter (caused by preemption
                                           or variable execution time)
    10. sleep(wake_up_time - current_time) // same as line 3
end_loop
```

Summary: Every line that manipulates time or performs I/O has problems.

Solution Approach: When I/O is Short (1)


- To solve the drift problem, use a periodic, hardware based timer to kick start each instance of periodic tasks. The tasks will be ready at the correct time instants.
- To minimize the jitter problem, perform the I/O in the timer interrupt handlers.
- Inside the handler, the processor will only be interrupted by another interrupt with a higher priority. See nested interrupts from previous lecture.
- As long as each task finishes before its end of period, I/O can be done at nearly the regular instants of timer interrupt, the highest regularity.
- The I/Os will still interfere with each other. But when the I/O times are short, this method works just fine. Also, you may use condition variables which will simplify the actual implementation.

Solution Approach: When I/O is Short (2)


- Initialization: // Lock mutex *wake_task_1_up*.
- Timer_interrupt_handler() // by convention, handlers execute before application tasks
 - {
 - // Do I/O ONLY – Why can't we do both work and I/O in the handler?
 - // Unlock mutex *wake_task_1_up*
 - }
- Task_1()
 - {
 - Loop
 - // Lock mutex *wake_task_1_up*
 - // Do non-I/O computation
 - End loop
 - }

Solution Approach: When I/O is Short (3)

Doing I/O in task body



Doing I/O inside timer ISR




Solutions Approach: When RT I/O is Too Long (1)


- In this case, hardware solution is called for. In high end A/D – D/A cards, there is a hardware timer for each input and output channel. For example,
 - ch 0 is input.
 - We can program its timer to sample the input line every 20 ms, starting at time 0.
 - put the data in ch 0's input buffer. (jitter free read)
 - Generate an interrupt to trigger Task_1.
 - Task_1 ISR: kickoff the Task_1 body.
 - Task_1 body: read the data, do work and put result to ch1 buffer.
 - ch 1 is output. We program its timer to output data from its buffer every 20 ms starting at 20 ms. (jitter_free output. Do not start at time zero, why?)
- As long as Task_1 body finishes its work within 20 ms, there is no jitter even if Task_1 body is preempted, why?

Solutions Approach: When RT I/O is Too Long (2)

Doing I/O inside timer ISR



Doing I/O with external hardware I/O support

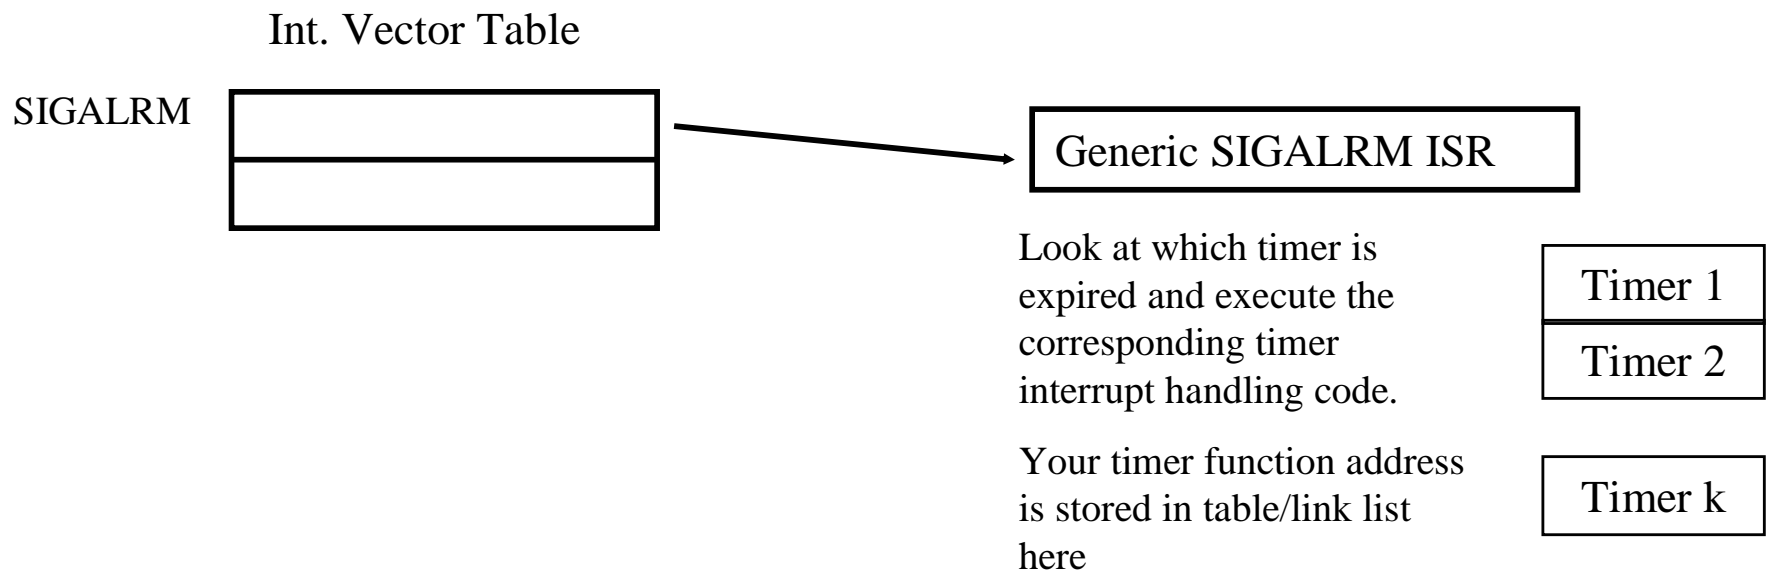


Signal Architecture (1)

- POSIX signals are implemented by software interrupts (e.g., SIGALRM)
 - Each periodic task thread will use a timer interrupt.
 - There can be as many threads as memory allows.
- How many real time clocks do we need? Why?
- How many entries in the interrupt vector table do we need?

Signal Architecture (2)

- A simple design: put the address of your ISR into the interrupt vector table. This is a valid and simple design. The drawback is that you need one alarm signal per thread. This can be done but not very efficient.
- An improved design. Recall that each thread has its own software timer. And all the threads wait on the same SIGALRM. To support the sharing of SIGALRM:



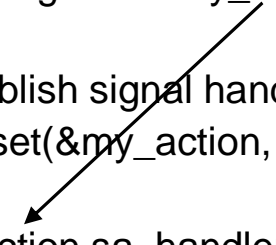
POSIX RT Timer Interrupt Handler

- POSIX Timer generates signal (software interrupts), SIGALRM. Action (ISR) for SIGALRM is therefore needed.
- Recall that you first define an interrupt type number which stores the address of your interrupt service routine.
- By POSIX coding convention, you define a generic action and then bind the action to a specific handler your write.

```
...
struct sigaction my_action, old_action //actions to catch a given signal arrives
...
// establish signal handler for SIGALRM
memset(&my_action, 0, sizeof(my_action)); // clear up the memory, a convention

my_action.sa_handler = timer_handler; // the "action" is to handle timer interrupt
//your ISR address for timer interrupt is copied into the sa_handler field

return_code = sigaction(SIGALRM, &my_action, &old_action); //binding action & signal
```



Summary

- Today
 - We have reviewed the issues in programming periodic tasks.
 - jitters and drifts.
 - How to use timers, and signal handler to construct a real time periodic tasks.
 - You are now ready to do Lab 2.
- Next Time
 - Basic signal processing.

Appendix: POSIX Real Time Clocks

- In POSIX RT defines the structure of time representation. There is at least 1 real time clock.
- Clock resolution is hardware dependent. (10 msec in PC without add-on high resolution real time clock cards)

- `#include time.h`

- ...

```
struct timespec current_time, clock_resolution
return_code = clock_gettime(CLOCK_REALTIME, &current_time);
//check return_code...
printf("%d %d current time in CLOCK_REALTIME is \n",
        current_time.tv_sec,           // the second portion
        current_time.tv_nsec);        // the fractional portion in nano-sec presentation
```

```
return_code = clock_getres(CLOCK_REALTIME, &clock_resolution)
//check return_code...
printf("%d %d CLOCK_REALTIME's resolution is \n",
        clock_resolution.tv_sec, clock_resolution.tv_nsec);
```

Appendix: Interval Timer Structure

POSIX timers are used to generate SOFTWARE INTERRUPTS (signals)

```
#include <signal.h>
```

```
#include <time.h>
```

```
//CREATE THE TIMER
```

```
timer_t timer_id;
```

```
//under POSIX, each thread defines its own timer for itself.
```

```
return_code= timer_create(CLOCK_REALTIME, NULL, &timer_id);
```

```
//NULL: no signal mask used to block other signals, if any, sent to this task
```

```
-----  
//INITIALIZE THE TIMER
```

```
struct itimerspec timer_spec, old_timer_spec;
```

```
//old timer allows us to save the old timer definition so you restore it if you need to
```

```
timer_spec.it_value.tv_sec = 10; // 1st expiration time
```

```
timer_spec.it_value.tv_nsec = 0;
```

```
timer_spec.it_interval.tv_sec = 0; // task period, 50Hz
```

```
timer_spec.it_interval.tv_nsec = 20000000;
```

```
timer_settime(timer_id, 0, &timer_spec, &old_timer_spec) //initialize the periodic timer
```

Appendix: Putting It Together

```
#include header files signal.h, time.h, errno.h, stdio.h //errno.h allows to decode the return code
```

```
void timer_handler( ) // it is invoked by the timer, not called by your software
{do your device I/O} // use global variables to communicate between main and handler
```

```
void main ()
{
    //ask user to input the task rate, max volt and min volt. Check for validity. If the resolution is too high
    //and if the voltages are too high/low
    //create your timer and initialize it
    //sets up your SIGALRM handler

    while (1) {
        sigpause(SIGALRM); //wait for signal SIGALRM
        //do your computation, make the handler code as small as possible to reduce jitter.
    }
}
```

Appendix: Using Signal Mask

Sigpause() is an old fashioned Unix function. POSIX RT defines sigsuspend() that uses a flexible signal mask. sigpause(signal_number) is still supported for backward compatibility.

```
void main ()
{
    sigset_t wait_signals, old_signals; //set of signals to wait for
    ...
    sigemptyset(&wait_signals); //initialize wait_signals to be an empty set
    sigaddset(&wait_signals, SIGALRM); //add SIGALRM to the set
    sigprocmask(SIG_BLOCK, &wait_signals, &old_signals); //install the signal set as a
    blocking mask
    while (1) {
        return_code = sigsuspend(&wait_signals); //wait for signal SIGALRM
        //do computation, make the handler code as small as possible to reduce jitter.
    }
}
```